

```
1 using Pkg; Pkg.activate() # since we need the extend-DR branch
```

```
Activating project at '~/.julia/environments/v1.10'
```



```
1 using Manifolds, Manopt, ManoptExamples
```

```
1 using NamedColors, Plots, BenchmarkTools
```

```
1 begin
2 paul_tol = load_paul_tol()
3 indigo = paul_tol["mutedindigo"]
4 green = paul_tol["mutedgreen"]
5 sand = paul_tol["mutedsand"]
6 olive = paul_tol["mutedolive"]
7 teal = paul_tol["mutedteal"]
8 wine = paul_tol["mutedwine"]
9 grey = paul_tol["mutedgrey"]
10 end;
```

Numerical Example for the acceleration and inertia on DR

The Rosenbrock problem

```
E = Euclidean(2; field=ℝ)
```

```
1 E = Euclidean(2)
```

```
M = MetricManifold(Euclidean(2; field=ℝ), ManoptExamples.RosenbrockMetric())
```

```
1 M = MetricManifold(E, ManoptExamples.RosenbrockMetric())
```

```
1 a = 2*10^5; b = 1; p0 = [0.1, 0.2]; p_star = [b, b^2];
```

```
f = RosenbrockCost(200000, 1)
```

```
1 f = ManoptExamples.RosenbrockCost(a, b)
```

```
∇f!! = RosenbrockGradient!!(200000, 1)
```

```
1 ∇f!! = ManoptExamples.RosenbrockGradient!!(a=a, b=b)
```

prox_f2 (generic function with 1 method)

```
1 begin
2 # These are described as "being defined" in the paper – I would prefer if that
  would be derived.
3 # f1, f2 hgere refers to the two summands of Rosenbrock, i.e.  $f = f_1 + f_2$ 
4 # (a) in-place variants
5 function prox_f1!(M, q, λ, p)
6     q .= [p[1], (p[2] + 2*a*λ*p[1]^2) / (1+2*a*λ)]
7     return q
8 end
9 function prox_f2!(M, q, λ, p)
10    q .= [
11        (p[1] + 2*λ*b) / (1+2*λ),
12        p[2] - ( 4*λ*(p[1] + 2*λ*b) * (p[1] - b) + 4 * λ^2 * (p[1]-b)^2 ) /
13            (1+2 * a * λ)^2
14    ]
15    return q
16 end
17 prox_f1(M, λ, p) = prox_f1!(M, copy(M, p), λ, p)
18 prox_f2(M, λ, p) = prox_f2!(M, copy(M, p), λ, p)
19 # The reflections here will work automatically, but I _hate_ phrasings like
  "can be easilu seen"
```

```
sc = StopWhenChangeLess(1.0e-14)
    |Δp| < 1.0e-14: not reached
```

```
1 sc = StopWhenChangeLess(1e-14)
```

Classical Douglas Rachford

```

s1 = # Solver state for `Manopt.jl`'s Douglas Rachford Algorithm
      After 32 iterations

      using an in place reflection.

      ## Stopping Criterion
      | $\Delta p$ | < 1.0e-14: reached
      This indicates convergence: Yes

      ## Debug
      :Stop = :Stop
      :All = [(:Iteration, "# %-6d"), (:Cost, "f(x): %f"), "\n", 25]

      ## Record
      (Iteration = RecordGroup([RecordIteration(), RecordCost()])),

```

```

1 # A first simple run
2 s1 = DouglasRachford(M, f, [prox_f1!, prox_f2!], p0;
3      $\alpha$  = i -> 0.5,
4      $\lambda$  = i -> 1.0,
5     debug = [ :Iteration, :Cost, "\n" , 25, :Stop],
6     record = [ :Iteration, :Cost],
7     stopping_criterion=sc,
8     evaluation = InplaceEvaluation(),
9     reflection_evaluation = InplaceEvaluation(),
10    return_state=true
11 )

```


```

Initial f(x): 7220.810000
# 25    f(x): 0.000000
The algorithm performed a step with a change (7.835532217989665e-15) less t
han 1.0e-14.

```

BenchmarkTools.Trial: 10000 samples with 1 evaluation.

Range (min ... max):	65.750 μ s ... 4.629 ms	GC (min ... max):	0.00% ... 96.67%
Time (median):	72.041 μ s	GC (median):	0.00%
Time (mean \pm σ):	76.446 μ s \pm 114.521 μ s	GC (mean \pm σ):	4.38% \pm 2.89%



Memory estimate: 52.59 KiB, allocs estimate: 1757.

```

1 # Benchmark
2 @benchmark DouglasRachford($M, $f, [$prox_f1!, $prox_f2!], $p0;
3      $\alpha$  = $(i -> 0.5),
4      $\lambda$  = $(i -> 1.0),
5     stopping_criterion = $(sc),
6     evaluation = $(InplaceEvaluation()),
7     reflection_evaluation = $(InplaceEvaluation()),
8 )

```

```
q1 = [1.0, 1.0]
```

```
1 q1 = get_solver_result(s1)
```

3.194886863360324e-24

```
1 f(M, q1)
```

Inertia

```
s2 = # Solver state for `Manopt.jl`'s Douglas Rachford Algorithm
      After 23 iterations

      using an in place reflection.

      ## Stopping Criterion
      |Δp| < 1.0e-14: reached
      This indicates convergence: Yes

      ## Debug
      :Stop = :Stop
      :All = [(:Iteration, "# %-6d"), (:Cost, "f(x): %f"), "\n", 10]

      ## Record
      (Iteration = RecordGroup([RecordIteration(), RecordCost()])),
```

```
1 s2 = DouglasRachford(M, f, [prox_f1!, prox_f2!], p0;
2   α = i -> 0.5,
3   λ = i -> 1.0,
4   θ = i -> 0.12, # Inertia
5   debug = [ :Iteration, :Cost, "\n" , 10, :Stop],
6   record = [ :Iteration, :Cost],
7   stopping_criterion=sc,
8   evaluation = InplaceEvaluation(),
9   reflection_evaluation = InplaceEvaluation(),
10  return_state=true
11 )
```

```
Initial f(x): 7220.810000
# 10    f(x): 0.000000
# 20    f(x): 0.000000
The algorithm performed a step with a change (3.815356980893987e-15) less than 1.0e-14.
```

```
q2 = [1.0, 1.0]
```

```
1 q2 = get_solver_result(s2)
```

3.549874073494553e-25

```
1 f(M, q2)
```

BenchmarkTools.Trial: 10000 samples with 1 evaluation.

Range (min ... max):	49.875 μ s ... 4.103 ms	GC (min ... max):	0.00% ... 97.14%
Time (median):	53.167 μ s	GC (median):	0.00%
Time (mean \pm σ):	56.115 μ s \pm 93.080 μ s	GC (mean \pm σ):	3.99% \pm 2.38%



Memory estimate: 38.98 KiB, allocs estimate: 1280.

```
1 @benchmark DouglasRachford($M, $f, [$prox_f1!, $prox_f2!], $p0;
2      $\alpha$  = $(i -> 0.5),
3      $\lambda$  = $(i -> 1.0),
4      $\theta$  = $(i -> 0.12),
5     stopping_criterion=$sc,
6     evaluation = $(InplaceEvaluation()),
7     reflection_evaluation = $(InplaceEvaluation()),
8 )
```

Acceleration

`s3` = # Solver state for 'Manopt.jl's Douglas Rachford Algorithm
After 10 iterations

using an in place reflection.

Stopping Criterion
| Δp | < 1.0e-14: reached
This indicates convergence: Yes

Debug
:Stop = :Stop
:All = [(:Iteration, "# %-6d"), (:Cost, "f(x): %f"), "\n", 10]

Record
(Iteration = RecordGroup([RecordIteration(), RecordCost()])),

```
1 s3 = DouglasRachford(M, f, [prox_f1!, prox_f2!], p0;
2      $\alpha$  = i -> 0.5,
3      $\lambda$  = i -> 1.0,
4     n = 3,
5     debug = [(:Iteration, :Cost, "\n" , 10, :Stop],
6     record = [ :Iteration, :Cost],
7     stopping_criterion=sc,
8     evaluation = InplaceEvaluation(),
9     reflection_evaluation = InplaceEvaluation(),
10    return_state=true
11 )
```

Initial $f(x)$: 7220.810000
10 $f(x)$: 0.000000
The algorithm performed a step with a change (2.2232198742534223e-15) less than 1.0e-14.



```
q3 = [1.0, 1.0]
```


```
1 q3 = get_solver_result(s3)
```

9.860761315262648e-27

```
1 f(M, q3)
```

BenchmarkTools.Trial: 10000 samples with 1 evaluation.

Range (min ... max):	63.875 μ s ... 3.374 ms	GC (min ... max):	0.00% ... 96.68%
Time (median):	66.292 μ s	GC (median):	0.00%
Time (mean \pm σ):	69.229 μ s \pm 82.043 μ s	GC (mean \pm σ):	3.08% \pm 2.55%



Memory estimate: 44.15 KiB, allocs estimate: 1842.

```
1 @benchmark DouglasRachford($M, $f, [$prox_f1!, $prox_f2!], $p0;  
2      $\alpha$  = $(i -> 0.5),  
3      $\lambda$  = $(i -> 1.0),  
4     n = 2,  
5     stopping_criterion=$sc,  
6     evaluation = $(InplaceEvaluation()),  
7     reflection_evaluation = $(InplaceEvaluation()),  
8 )
```

Acceleration *and* Interatia

```

s4 = # Solver state for 'Manopt.jl's Douglas Rachford Algorithm
      After 13 iterations

      using an in place reflection.

      ## Stopping Criterion
      |Δp| < 1.0e-14: reached
      This indicates convergence: Yes

      ## Debug
      :Stop = :Stop
      :All = [(:Iteration, "# %-6d"), (:Cost, "f(x): %f"), "\n", 10]

      ## Record
      (Iteration = RecordGroup([RecordIteration(), RecordCost()])),

```


```

1 s4 = DouglasRachford(M, f, [prox_f1!, prox_f2!], p0;
2   α = i -> 0.5,
3   λ = i -> 1.0,
4   θ = i -> 0.12, # Inertia
5   n = 3,
6   debug = [ :Iteration, :Cost, "\n" , 10, :Stop],
7   record = [ :Iteration, :Cost],
8   stopping_criterion=sc,
9   evaluation = InplaceEvaluation(),
10  reflection_evaluation = InplaceEvaluation(),
11  return_state=true
12 )

```

Initial f(x): 7220.810000
 # 10 f(x): 0.000000
 The algorithm performed a step with a change (8.560244302824867e-15) less than 1.0e-14. ?

BenchmarkTools.Trial: 10000 samples with 1 evaluation.
 Range (min ... max): 82.666 μs ... 3.397 ms | GC (min ... max): 0.00% ... 96.09%
 Time (median): 84.958 μs | GC (median): 0.00%
 Time (mean ± σ): 88.666 μs ± 92.404 μs | GC (mean ± σ): 3.07% ± 2.87%



Memory estimate: 54.32 KiB, allocs estimate: 2388.

```

1 @benchmark DouglasRachford($M, $f, [$prox_f1!, $prox_f2!], $p0;
2   α = $(i -> 0.5),
3   λ = $(i -> 1.0),
4   θ = i -> 0.12, # Inertia
5   n = 3,
6   stopping_criterion=$sc,
7   evaluation = $(InplaceEvaluation()),
8   reflection_evaluation = $(InplaceEvaluation()),
9 )

```

Quasi Newton

A sanity check – quasi Newton should be able to do this reasonably quick as well.

```
s5 =  
# Solver state for `Manopt.jl`'s Quasi Newton Method  
After 213 iterations  
  
## Parameters  
* direction update:          limited memory Manopt.InverseBFGS (size 5), projections, an  
* retraction method:        ManifoldsBase.ExponentialRetraction()  
* vector transport method: ParallelTransport()  
  
## Stepsize  
WolfePowellLinesearch(DefaultManifold(), 0.0001, 0.999) with keyword arguments  
  * retraction_method = ManifoldsBase.ExponentialRetraction()  
  * vector_transport_method = ParallelTransport()  
  
## Stopping Criterion  
|Δp| < 1.0e-14: reached  
This indicates convergence: Yes  
  
## Debug  
  :Stop = :Stop  
  :All = [(:Iteration, "# %-6d"), (:Cost, "f(x): %f"), "\n", 100]  
  
## Record  
(Iteration = RecordGroup([RecordIteration(), RecordCost()])),
```

```
1 s5 = quasi_Newton(E, f, ∇f!!, p0;  
2   debug = [(:Iteration, :Cost, "\n" , 100, :Stop],  
3   record = [ :Iteration, :Cost],  
4   memory_size=5,  
5   stopping_criterion=sc,  
6   evaluation = InplaceEvaluation(),  
7   return_state=true  
8 )
```

```
Initial f(x): 7220.810000  
# 100    f(x): 0.097674  
# 200    f(x): 0.000001  
The algorithm performed a step with a change (2.9790409838967277e-15) less  
than 1.0e-14.
```

```
q5 = [1.0, 1.0]
```

```
1 q5 = get_solver_result(s5)
```

```
0.0
```


```
1 f(E, q5)
```



```
@benchmark s5 =
```

BenchmarkTools.Trial: 10000 samples with 1 evaluation.

Range (min ... max):	125.958 μ s ... 3.608 ms	GC (min ... max):	0.00% ... 95.11%
Time (median):	131.250 μ s	GC (median):	0.00%
Time (mean \pm σ):	137.208 μ s \pm 121.987 μ s	GC (mean \pm σ):	3.50% \pm 3.78%



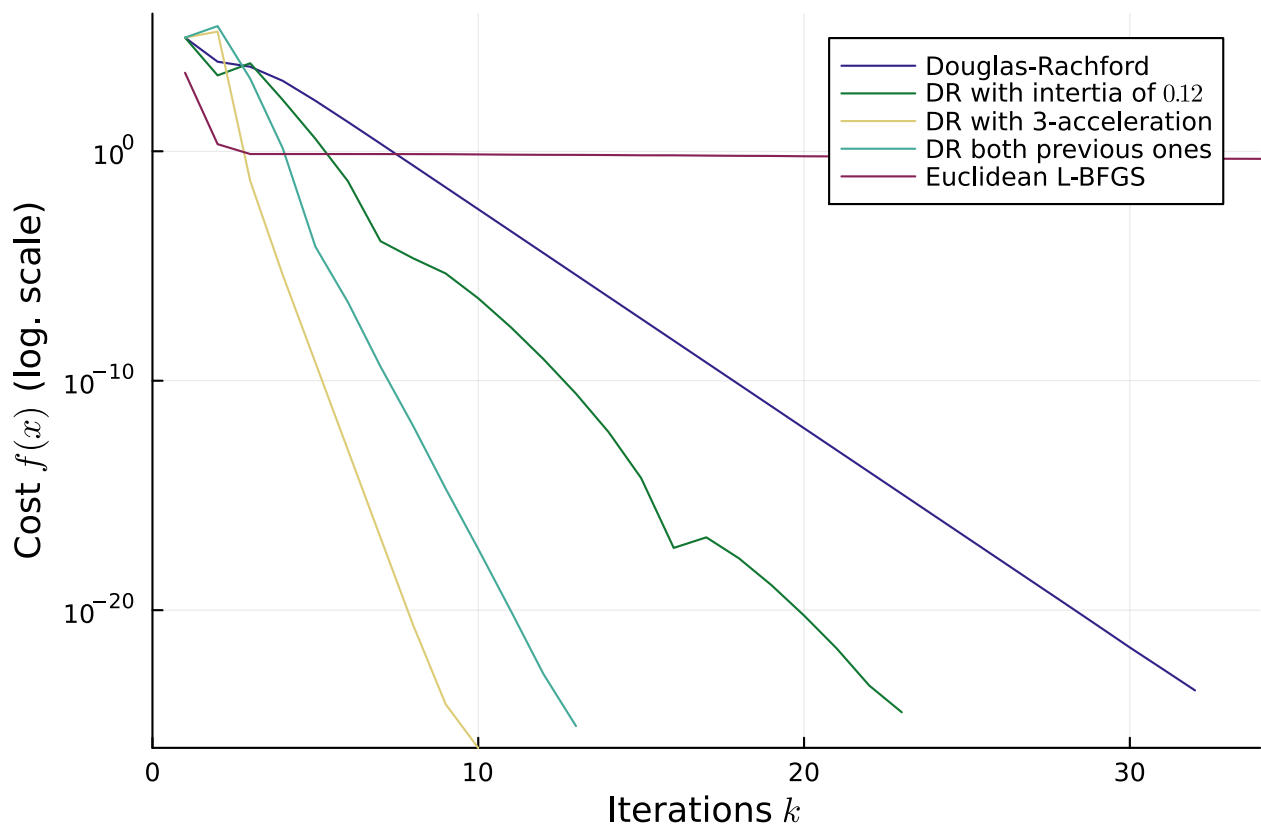
Memory estimate: 98.37 KiB, allocs estimate: 1123.

```
1 @benchmark s5 = quasi_Newton($E, $f, $Vf!!, $p0;  
2     memory_size=5,  
3     stopping_criterion=$sc,  
4     evaluation = $(InplaceEvaluation()),  
5 )
```

Summary

```
1 iterates = [ [e[1] for e in get_record(s, :Iteration)] for s in [s1,s2, s3, s4,  
s5]];
```

```
1 costs = [ [e[2] for e in get_record(s, :Iteration)] for s in [s1,s2, s3, s4,  
s5]];
```




```

1 begin
2     fig = plot(
3         xlabel=raw"Iterations $k$", ylabel=raw"Cost $f(x)$ (log. scale)",
4         axis=:log,
5         ylim = (1e-26,1e6),
6         xlim = (0,34),
7     );
8     plot!(fig, iterates[1], costs[1], color=indigo, label="Douglas-Rachford");
9     plot!(fig, iterates[2], costs[2], color=green, label=raw"DR with inertia of
10    $0.12$");
11    plot!(fig, iterates[3], costs[3], color=sand, label=raw"DR with 3-
12    acceleration");
13    plot!(fig, iterates[4], costs[4], color=teal, label=raw"DR both previous
14    ones");
15    fig
16    plot!(fig, iterates[5], costs[5], color=wine, label=raw"Euclidean L-BFGS");
17    fig
18 end

```

When we look at Quasi Newton, it reaches even numerical zero, but only relatively late:



```

1 begin
2     fig2 = deepcopy(fig)
3     plot!(fig2, xlim=(0,214))
4 end

```